



Recap: Execution of SQL Queries

```
select dept_name, avg(salary)
from instructor join department
where budget > 100000
group by dept_name
having avg (salary) > 42000;
```

is equivalent to RA expression

$$\prod_{dept_name, avs} (\sigma_{avs > 42000} (\pi_{dept_name} G avg(salary) as avs
(\sigma_{budget > 100000} (instructor \bowtie department))))$$



Recap: Execution of SQL Queries

```
select dept_name, avg(salary)
from instructor join department
where budget > 100000
group by dept_name
having avg (salary) > 42000;
```

Is executed in the following steps:

- **from**: do the cross product or join (in his case) of the tables in the **from** clause.
- **where**: apply the condition in the **where** clause to every tuple of the table produced by the **from** clause.
- **group by**: aggregate and reorg table into a new table containing only the **group by** attributes and any aggregates referenced in the **having** or **select** clauses.
- **having**: apply condition in the **having** clause to every tuple of the table produced by the **group by**.
- **select**: project onto the requested attributes



Notes on Execution of SQL Queries

- A SQL query produces a table as output
- In fact, every of the steps on previous slide starts with one or more table, and produces a table
- Thus, the meaning of nesting in the **from** clause is obvious
- For nesting in the **where**, **having**, and **select** clauses, keep in mind that these clauses are applied to each row of the table entering it
- So if there are correlation variables, the subqueries may have to be repeatedly executed
- Exception for **group by** in many DBMS: if you group by a primary key, other attributes in that table can be referred to after group-by:

```
select D.dept_name, D.budget, avg(salary)  
from instructor I join department D  
where D.budget > 100000  
group by D.dept_name
```



Chapter 4: Intermediate SQL

- Join Expressions
- Views
- Transactions
- Integrity Constraints
- SQL Data Types and Schemas
- Authorization



Joined Relations

- **Join operations** take two relations and return as a result another relation.
- A join operation is a Cartesian product which requires that tuples in the two relations match (under some condition). It also specifies the attributes that are present in the result of the join.
- The join operations are typically used as subquery expressions in the **from** clause.



Join operations – Example

■ Relation *course*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

■ Relation *prereq*

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

- Note: prereq information missing for CS-315 and course information missing for CS-437.



Outer Join

- An extension of the join operation that avoids loss of information.
- Computes the join and then adds tuples from one relation that does not match tuples in the other relation to the result of the join.
- Uses *null* values.

Left Outer Join

- *course natural left outer join prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prere_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>

Note: read *prere_id* as *prereq_id*



Right Outer Join

- course natural right outer join prereq

course_id	title	dept_name	credits	prere_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-347	null	null	null	CS-101

Full Outer Join

- course natural full outer join prereq

course_id	title	dept_name	credits	prere_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	null
CS-347	null	null	null	CS-101



Join Types vs. Join Conditions

- **Join operations** take two relations and return as a result another relation.
- These additional operations are typically used as subquery expressions in the **from** clause
- **Join condition** – defines which tuples in the two relations match, and what attributes are present in the result of the join.
- **Join type** – defines how tuples in each relation that do not match any tuple in the other relation (based on the join condition) are treated.

<i>Join types</i>
inner join
left outer join
right outer join
full outer join

<i>Join Conditions</i>
natural
on <predicate>
using (A_1, A_1, \dots, A_n)



Joined Relations – Examples

- course **inner join** prereq **on**
 $course.course_id = prereq.course_id$

course_id	title	dept_name	credits	prere_id	course_id
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190

- course **left outer join** prereq **on**
 $course.course_id = prereq.course_id$

course_id	title	dept_name	credits	prere_id	course_id
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190
CS-315	Robotics	Comp. Sci.	3	null	null



Joined Relations – Examples

- course **natural right outer join** prereq

course_id	title	dept_name	credits	prere_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-347	null	null	null	CS-101

- course **full outer join** prereq using (course_id)

course_id	title	dept_name	credits	prere_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	null
CS-347	null	null	null	CS-101



Example

- Find the course ID, semester, year and title of each course offered by the Comp. Sci. department

```
select section.course_id, semester, year, title  
from section, course  
where section.course_id = course.course_id and  
      dept_name = 'Comp. Sci.'
```

- Same with join operation

```
select section.course_id, semester, year, title  
from section join course  
where dept_name = 'Comp. Sci.'
```



Views

- In some cases, it is not desirable for all users to see the entire logical model (that is, all the actual relations stored in the database.)
- Consider a person (or an application) who needs to know an instructors name and department, but not the salary. This person should see a relation described, in SQL, by

```
select ID, name, dept_name  
from instructor
```

- A **view** provides a mechanism to hide certain data from the view of certain users.
- Any relation that is not of the conceptual model but is made visible to a user as a “virtual relation” is called a **view**.



View Definition

- A view is defined using the **create view** statement which has the form

create view *v* as <query expression>

where <query expression> is any legal SQL expression.
The view name is represented by *v*.

- Once a view is defined, the view name can be used to refer to the virtual relation that the view generates.
- View definition is not the same as creating a new relation by evaluating the query expression
 - Rather, a view definition causes the saving of an expression; the expression is substituted into queries using the view.



Example Views

- A view of instructors without their salary

```
create view faculty as
```

```
select ID, name, dept_name  
from instructor
```

- Find all instructors in the Biology department

```
select name
```

```
from faculty
```

```
where dept_name = 'Biology'
```

- Create a view of department salary totals

```
create view departments_total_salary(dept_name, total_salary) as
```

```
select dept_name, sum (salary)
```

```
from instructor
```

```
group by dept_name;
```



Views Defined Using Other Views

- **create view *physics_fall_2009* as**
select course.course_id, sec_id, building, room_number
from course, section
where course.course_id = section.course_id
and course.dept_name = 'Physics'
and section.semester = 'Fall'
and section.year = '2009';

- **create view *physics_fall_2009_watson* as**
select course_id, room_number
from *physics_fall_2009*
where building= 'Watson';



View Expansion

- Expand use of a view in a query/another view

```
create view physics_fall_2009_watson as
  (select course_id, room_number
   from (select course.course_id, building,
              room_number
            from course, section
           where course.course_id = section.course_id
             and course.dept_name = ' Physics'
             and section.semester = ' Fall'
             and section.year = ' 2009')
    where building = ' Watson' ;
```



Views Defined Using Other Views

- One view may be used in the expression defining another view,
- A view relation v_1 is said to *depend directly* on a view relation v_2 if v_2 is used in the expression defining v_1
- A view relation v_1 is said to *depend on* view relation v_2 if either v_1 depends directly to v_2 or there is a path of dependencies from v_1 to v_2
- A view relation v is said to be *recursive* if it depends on itself.



View Expansion

- A way to define the meaning of views defined in terms of other views.
- Let view v_1 be defined by an expression e_1 that may itself contain uses of view relations.
- View expansion of an expression repeats the following replacement step:

repeat
 Find any view relation v_i in e_1
 Replace the view relation v_i by the expression defining v_i
until no more view relations are present in e_1
- As long as the view definitions are not recursive, this loop will terminate.



How Can your DBMS Maintain Views?

- If a view is often used in queries, it might be very inefficient to recompute the view from scratch every time it is used.



How Can your DBMS Maintain Views?

- If a view is often used in queries, it might be very inefficient to recompute the view from scratch every time it is used.
- To deal with this, your DBMS will often materialize and then maintain the view under updates to the underlying tables (“materialized views”)

- Example 1: the faculty view on slide 15:
 - Every insert, delete, update to instructor can be translated into a suitable insert, delete, or update on the materialized view



How Can your DBMS Maintain Views?

- If a view is often used in queries, it might be very inefficient to recompute the view from scratch every time it is used.
- To deal with this, your DBMS will often materialize and then maintain the view under updates to the underlying tables (“materialized views”)
- Example 1: the faculty view on slide 15:
 - Every insert, delete, update to instructor can be translated into a suitable insert, delete, or update on the materialized view
- Example 2: the departments_total_salary view on slide 15:
 - Every insert, delete, and update to instructor can be applied to the view by increasing, decreasing, or updating department salary totals
- The best strategy depends on number of view accesses versus number of updates to the underlying table



Update of a View

- Add a new tuple to *faculty* view which we defined earlier
 - `insert into faculty values ('30765', 'Green', 'Music');`
- This insertion could be represented by the insertion of the tuple
 - `('30765', 'Green', 'Music', null)`into the *instructor* relation
- Is this legal?



Update of a View

- Add a new tuple to *faculty* view which we defined earlier
 - `insert into faculty values ('30765', 'Green', 'Music');`
- This insertion could be represented by the insertion of the tuple
 - `('30765', 'Green', 'Music', null)`into the *instructor* relation
- Is this legal?
- Yes, assuming we may assign a null value to salary
- If salary attribute was NOT NULL, cannot do this



Some Updates cannot be Translated Uniquely

- `create view instructor_info as
 select ID, name, building
 from instructor, department
 where instructor.dept_name= department.dept_name;`
- `insert into instructor info values (' 69987' , ' White' , ' Taylor');`
 - ▶ which department, if multiple departments in Taylor?
 - ▶ what if no department is in Taylor?
- Most SQL implementations allow updates only on simple views
 - The **from** clause has only one database relation.
 - The **select** clause contains only attribute names of the relation, and does not have any expressions, aggregates, or **distinct** specification.
 - Any attribute not listed in the **select** clause can be set to null
 - The query does not have a **group by** or **having** clause.
- But e.g., Oracle allows some updates on views with several relations



And Some Not at All

- **create view** *history_instructors* **as**
select *
from *instructor*
where *dept_name*= 'History' ;
- Insert ('25566', 'Brown', 'Biology', 100000) into
history_instructors



Transactions

- Unit of work
- Atomic transaction
 - either fully executed or rolled back as if it never occurred
- Isolation from concurrent transactions
- Transactions begin implicitly
 - Ended by **commit work** or **rollback work**
- But default on most databases: each SQL statement commits automatically
 - Can turn off auto commit for a session (e.g. using API)
 - In SQL:1999, can use: **begin atomic end**



Integrity Constraints

- Integrity constraints guard against accidental damage to the database, by ensuring that authorized changes to the database do not result in a loss of data consistency.
 - A checking account must have a balance greater than \$10,000.00.
 - A salary of a bank employee must be at least \$4.00 an hour.
 - A customer must have a (non-null) phone number.



Constraints on a Single Relation

- **not null**
- **primary key**
- **unique**
- **check (P)**, where P is a predicate



Not Null and Unique Constraints

■ not null

- Declare *name* and *budget* to be **not null**

name varchar(20) not null

budget numeric(12,2) not null

■ unique (A_1, A_2, \dots, A_m)

- The unique specification states that the attributes A_1, A_2, \dots, A_m form a candidate key.
- Candidate keys are permitted to be null (in contrast to primary keys).



The check clause

■ **check (P)**

where P is a predicate

Example: ensure that semester is one of fall, winter, spring or summer:

```
create table section (
    course_id varchar (8),
    sec_id varchar (8),
    semester varchar (6),
    year numeric (4,0),
    building varchar (15),
    room_number varchar (7),
    time_slot_id varchar (4),
    primary key (course_id, sec_id, semester, year),
    check (semester in (' Fall', ' Winter', ' Spring',
    ' Summer'))
```



Referential Integrity

- Ensures that a value that appears in one relation for a given set of attributes also appears for a certain set of attributes in another relation.
 - Example: If “Biology” is a department name appearing in one of the tuples in the *instructor* relation, then there exists a tuple in the *department* relation for “Biology”.
- Let A be a set of attributes. Let R and S be two relations that contain attributes A and where A is the primary key of S. A is said to be a **foreign key** of R if for any values of A appearing in R these values also appear in S.



Cascading Actions in Referential Integrity

- **create table course (**
course_id char(5) primary key,
title varchar(20),
dept_name varchar(20) references department
)
- **create table course (**
...
dept_name varchar(20),
foreign key (dept_name) references department
on delete cascade
on update cascade,
...
)
- alternative actions to cascade: **set null, set default**



Complex Check Clauses

- **check** (*time_slot_id* in (select *time_slot_id* from *time_slot*))
 - why not use a foreign key here?
- Every section has at least one instructor teaching the section.
 - how to write this?
- Unfortunately: subquery in check clause not supported by pretty much any database
 - Alternative: triggers (later)
- **create assertion** <assertion-name> **check** <predicate>;
 - Also not supported by anyone



Built-in Data Types in SQL

- **date:** Dates, containing a (4 digit) year, month and date
 - Example: **date** ‘2005-7-27’
- **time:** Time of day, in hours, minutes and seconds.
 - Example: **time** ‘09:00:30’ **time** ‘09:00:30.75’
- **timestamp:** date plus time of day
 - Example: **timestamp** ‘2005-7-27 09:00:30.75’
- **interval:** period of time
 - Example: **interval** ‘1’ day
 - Subtracting a date/time/timestamp value from another gives an interval value
 - Interval values can be added to date/time/timestamp values



Other Features

- **create table** *student*
*(ID varchar (5),
name varchar (20) not null,
dept_name varchar (20),
tot_cred numeric (3,0) default 0,
primary key (ID))*
- **create index** *studentID* **index on** *student(ID)*
- Large objects
 - *book review* **clob(10KB)**
 - *image blob(10MB)*, *movie blob(2GB)*
- **create type** construct in SQL creates user-defined type

create type *Dollars as numeric (12,2) final*
 - **create table** *department*
*(dept_name varchar (20),
building varchar (15),
budget Dollars);*



- **create domain** construct in SQL-92 creates user-defined domain types

```
create domain person_name char(20) not null
```

- Types and domains are similar. Domains can have constraints, such as **not null**, specified on them.

- **create domain** *degree_level* **varchar**(10)

```
constraint degree_level_test
```

```
check (value in ('Bachelors', 'Masters', 'Doctorate'));
```

- Large objects (photos, videos, CAD files, etc.) are stored as a *large object*:

- **blob**: binary large object -- object is a large collection of uninterpreted binary data (whose interpretation is left to an application outside of the database system)
- **clob**: character large object -- object is a large collection of char data
- When a query returns a large object, a pointer is returned rather than the large object itself.



Authorization

Forms of authorization on parts of the database:

- **Read** - allows reading, but not modification of data.
- **Insert** - allows insertion of new data, but not modification of existing data.
- **Update** - allows modification, but not deletion of data.
- **Delete** - allows deletion of data.

Forms of authorization to modify the database schema

- **Index** - allows creation and deletion of indices.
- **Resources** - allows creation of new relations.
- **Alteration** - allows addition or deletion of attributes in a relation.
- **Drop** - allows deletion of relations.



Authorization Specification in SQL

- The **grant** statement is used to confer authorization

grant <privilege list>

on <relation name or view name> to <user list>

- <user list> is:

- a user-id
- **public**, which allows all valid users the privilege granted
- A role (more on this later)

- Granting a privilege on a view does not imply granting any privileges on the underlying relations.
- The grantor of the privilege must already hold the privilege on the specified item (or be the database administrator).



Privileges in SQL

- **select:** allows read access to relation, or the ability to query using the view
 - Example: grant users U_1 , U_2 , and U_3 **select** authorization on the *branch* relation:
$$\text{grant select on } \textit{instructor} \text{ to } U_1, U_2, U_3$$
- **insert:** the ability to insert tuples.
- **update:** the ability to update using the SQL update statement.
- **delete:** the ability to delete tuples.
- **all privileges:** used as a short form for all the allowable privileges.



Revoking Authorization in SQL

- The **revoke** statement is used to revoke authorization.

revoke <privilege list>

on <relation name or view name> from <user list>

- Example:

revoke select on branch from U_1, U_2, U_3

- <privilege-list> may be **all** to revoke all privileges the revoker may hold.
- If <revoker-list> includes **public**, all users lose the privilege except those granted it explicitly.
- If the same privilege was granted twice to the same user by different grantees, the user may retain the privilege after the revocation.
- All privileges that depend on the privilege being revoked are also revoked.



Roles

- **create role** *instructor*;
- Privileges can be granted to roles:
 - **grant select on** *takes* **to** *instructor*;
- Roles can be granted to users, as well as to other roles
 - **create role** *student*
 - **grant** *instructor* **to** Amit;
 - **create role** *dean*;
 - **grant** *instructor* **to** *dean*;
 - **grant** *dean* **to** Satoshi;



Authorization on Views

- **create view geo_instructor as**
(select *
from instructor
where dept_name = 'Geology');
- **grant select on geo_instructor to staff**
- Suppose that a staff member issues
 - **select ***
from geo_instructor;
- What if
 - staff does not have permissions on *instructor*?
 - creator of view did not have some permissions on *instructor*?